

Welcome to the Kappa Zoo

David A. Eubanks

Table of contents

1	Welcome	1
1.1	Table of Contents	2

1 Welcome

This site—a work in progress—explains rater reliability and how to measure it. Much of the content is new, using a general model to explain the features of existing methods from the measurement literature, in the context of more modern approaches from the machine learning (ML) literature. The sythesis of these two fields is the t-a-p model, which combines the intuitiveness of the measurement ideas with the statistical power of the ML algorithms.

Use the navigation bar at the top of the page to explore the chapters, or jump directly to specific sections using the links below. If you are new to t-a-p models, I recommend starting with the [Chapter 1](#). Then if you want to dive into theory, the subsequent two chapters lay out the statistical derivations and relationship to older methods, for example showing how the Fleiss kappa is a special case of a t-a-p model.

If you'd rather jump straight to applications, you can skip the theory and go to the [The t-a-p App](#) page, which is a user manual for the R Shiny app that accompanies this site. With it you can simulate data or load your own, then model the data for rater accuracy.

The Kappa Zoo includes a growing collection of real-world data sets and the model parameters estimated from them.

1.1 Table of Contents

- [Chapter 1: Quantifying Agreement](#)
- [Chapter 2: Estimating t-a-p Parameters](#)
- [Chapter 3: Kappa Statistics](#)
- [Chapter 4: The t-a-p App](#)
- [Chapter 5: The t-a-p R Package](#)
- [The Kappa Zoo](#)
- [References](#)
- [Appendix A: Statistical Details](#)

1.1.1 Getting Started

I recommend starting with the [Introduction](#). Even if you're familiar with the concept of rater reliability, you'll want to understand how t-a-p models work, and the introduction gives an overview.

For the technical details of the t-a-p model, see [Chapter 1](#).